




# Search: dynamic programming



# Dynamic programming



Minimum cost path from state  $s$  to a end state:


$$\text{FutureCost}(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ \min_{a \in \text{Actions}(s)} [\text{Cost}(s, a) + \text{FutureCost}(\text{Succ}(s, a))] & \text{otherwise} \end{cases}$$

# Motivating task



## Example: route finding


Find the minimum cost path from city 1 to city  $n$ , only moving forward. It costs  $c_{ij}$  to go from  $i$  to  $j$ .



Observation: future costs only depend on current city

# Dynamic programming

**State:** ~~past sequence of actions~~ current city



**Exponential saving in time and space!**

# Dynamic programming



## Algorithm: dynamic programming

```
def DynamicProgramming( $s$ ):  
    If already computed for  $s$ , return cached answer.  
    If IsEnd( $s$ ): return solution  
    For each action  $a \in \text{Actions}(s)$ : ...
```

[semi-live solution: Dynamic Programming]



## Assumption: acyclicity

The state graph defined by  $\text{Actions}(s)$  and  $\text{Succ}(s, a)$  is acyclic.

# Dynamic programming




## Key idea: state

A **state** is a summary of all the past actions sufficient to choose future actions **optimally**.

past actions (all cities)      1 3 4 6

state (current city)          1 3 4 6

# Handling additional constraints

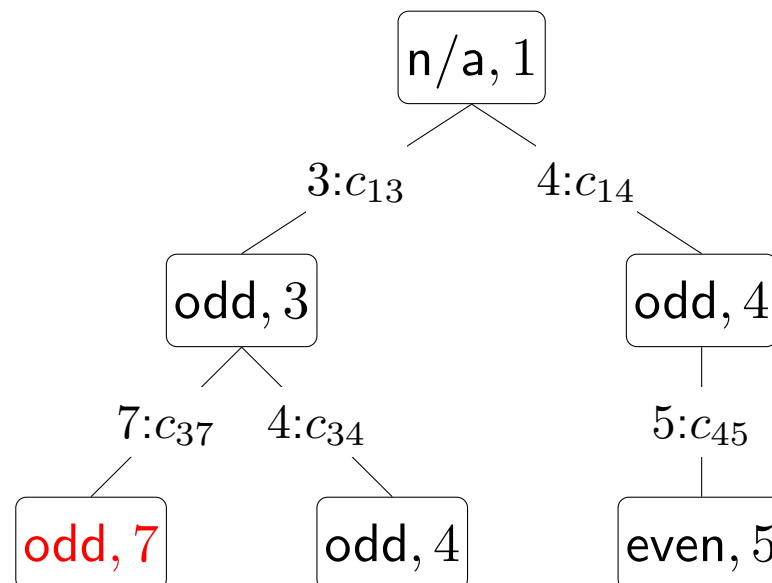


## Example: route finding

Find the minimum cost path from city 1 to city  $n$ , only moving forward. It costs  $c_{ij}$  to go from  $i$  to  $j$ .

**Constraint: Can't visit three odd cities in a row.**

**State:** (whether previous city was odd, current city)





answer in chat

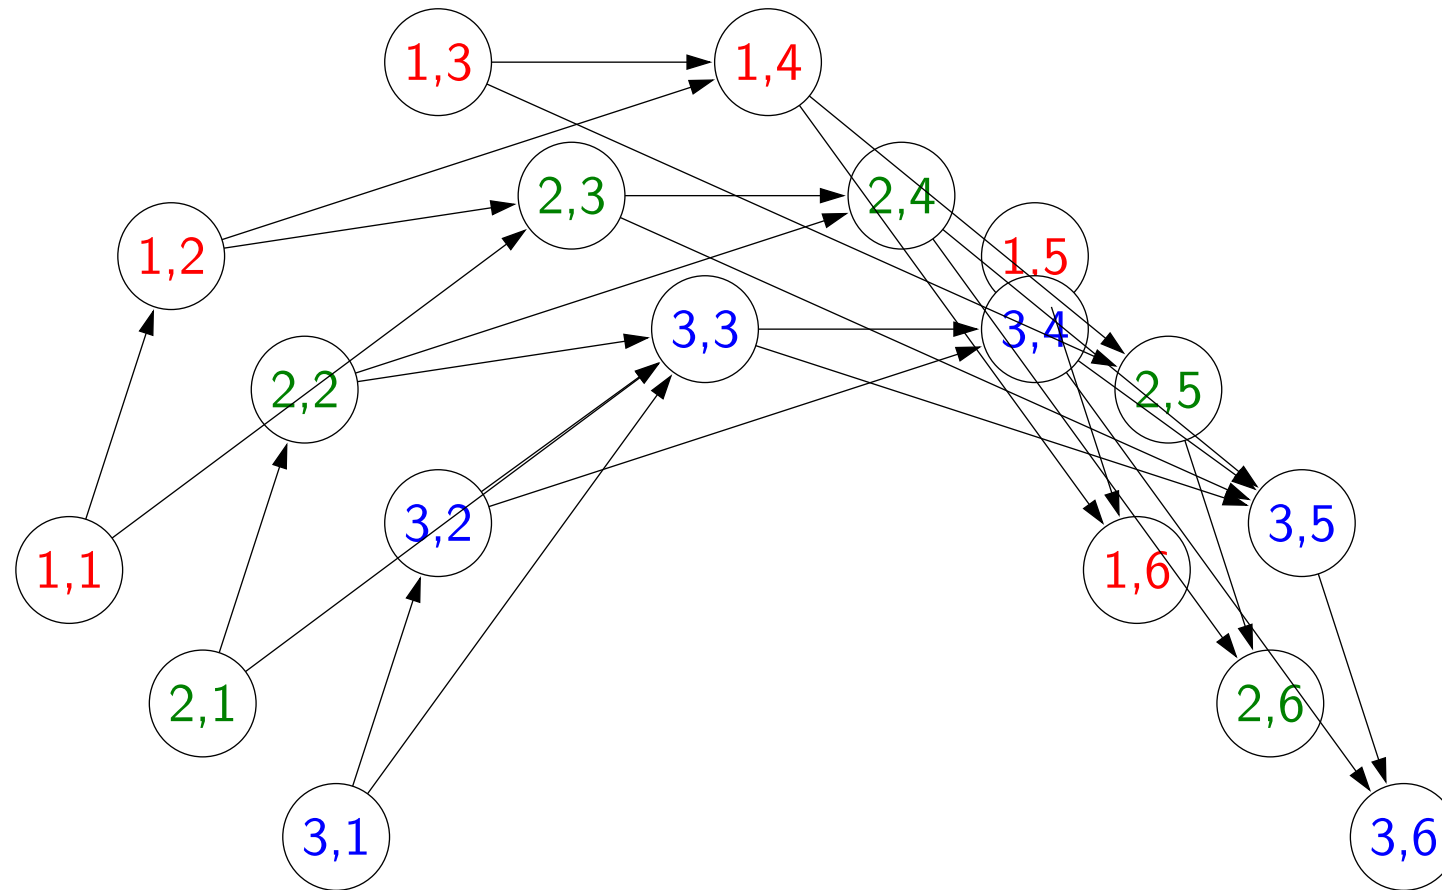
# Question

Objective: travel from city 1 to city  $n$ , visiting at least 3 odd cities. What is the minimal state?



# State graph

State: (min(number of odd cities visited, 3), current city)





answer in chat

# Question

Objective: travel from city 1 to city  $n$ , visiting more odd than even cities. What is the minimal state?




# Summary

- **State**: summary of past actions sufficient to choose future actions optimally
- **Dynamic programming**: backtracking search with **memoization** — potentially exponential savings

Dynamic programming only works for acyclic graphs...what if there are cycles?

# Dynamic Programming Review



$$\text{FutureCost}(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ \min_{a \in \text{Actions}(s)} [\text{Cost}(s, a) + \text{FutureCost}(\text{Succ}(s, a))] & \text{otherwise} \end{cases}$$



**Key idea: state**

A **state** is a summary of all the past actions sufficient to choose future actions **optimally**.




# Search: uniform cost search



# Ordering the states

**Observation:** prefixes of optimal path are optimal



**Key:** if graph is acyclic, dynamic programming makes sure we compute  $\text{PastCost}(s)$  before  $\text{PastCost}(s')$

If graph is cyclic, then we need another mechanism to order states...

# Uniform cost search (UCS)



**Key idea: state ordering**


UCS enumerates states in order of increasing past cost.




**Assumption: non-negativity**

All action costs are non-negative:  $\text{Cost}(s, a) \geq 0$ .

UCS in action:



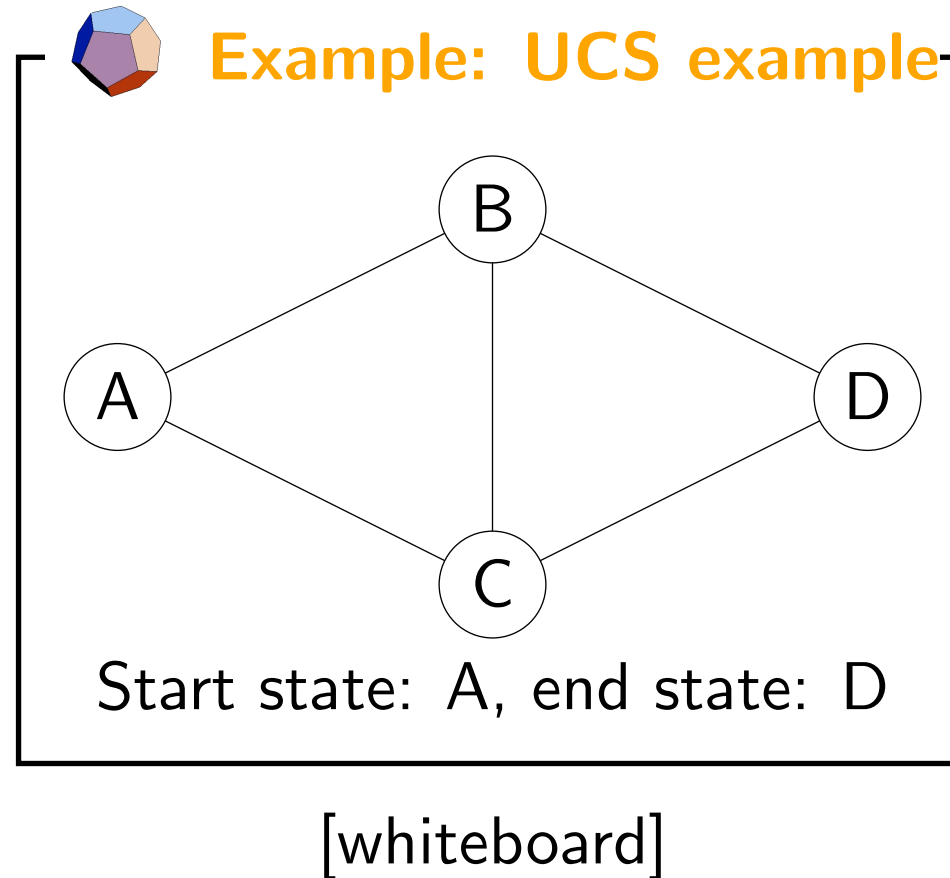
# High-level strategy



- **Explored**: states we've found the optimal path to
- **Frontier**: states we've seen, still figuring out how to get there cheaply
- **Unexplored**: states we haven't seen



# Uniform cost search example



Minimum cost path:

$A \rightarrow B \rightarrow C \rightarrow D$  with cost 3




Search:  $A^*$





# A\* algorithm


UCS in action:



A\* in action:



# Can uniform cost search be improved?




**Problem:** UCS orders states by cost from  $s_{\text{start}}$  to  $s$

**Goal:** take into account cost from  $s$  to  $s_{\text{end}}$


# Exploring states

UCS: explore states in order of  $\text{PastCost}(s)$



Ideal: explore in order of  $\text{PastCost}(s) + \text{FutureCost}(s)$

A\*: explore in order of  $\text{PastCost}(s) + h(s)$



## Definition: Heuristic function

A heuristic  $h(s)$  is any estimate of  $\text{FutureCost}(s)$ .

# A\* search




**Algorithm: A\* search [Hart/Nilsson/Raphael, 1968]**

Run uniform cost search with **modified edge costs**:

$$\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s)$$

**Intuition:** add a penalty for how much action  $a$  takes us away from the end state

**Example:**




$$\text{Cost}'(C, B) = \text{Cost}(C, B) + h(B) - h(C) = 1 + (3 - 2) = 2$$

# An example heuristic

Will any heuristic work?

No.

Counterexample:



Doesn't work because of **negative modified edge costs!**



# Consistent heuristics




## Definition: consistency

A heuristic  $h$  is **consistent** if


- $\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$
- $h(s_{\text{end}}) = 0$ .

Condition 1: needed for UCS to work (triangle inequality).



Condition 2:  $\text{FutureCost}(s_{\text{end}}) = 0$  so match it.

# Correctness of $A^*$




## **Proposition: correctness**

If  $h$  is consistent,  $A^*$  returns the minimum cost path.

# Proof of A\* correctness

- Consider any path  $[s_0, a_1, s_1, \dots, a_L, s_L]$ :




- Key identity:

$$\underbrace{\sum_{i=1}^L \text{Cost}'(s_{i-1}, a_i)}_{\text{modified path cost}} = \underbrace{\sum_{i=1}^L \text{Cost}(s_{i-1}, a_i)}_{\text{original path cost}} + \underbrace{h(s_L) - h(s_0)}_{\text{constant}}$$

- Therefore, A\* (finding the minimum cost path using modified costs) solves the original problem (even though edge costs are all different!)

# Efficiency of A\*



## Theorem: efficiency of A\*

A\* explores all states  $s$  satisfying  
 $\text{PastCost}(s) \leq \text{PastCost}(s_{\text{end}}) - h(s)$

Interpretation: the larger  $h(s)$ , the better

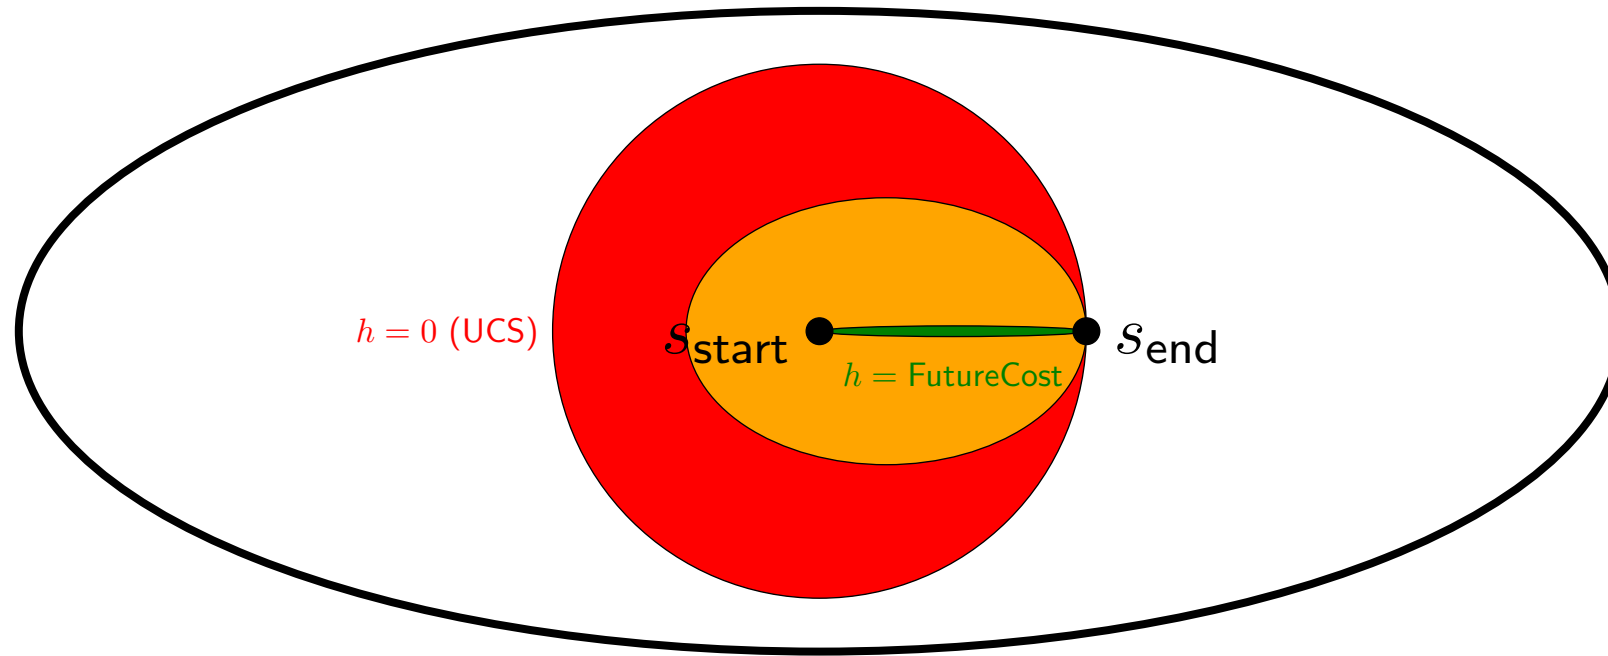
Proof: A\* explores all  $s$  such that

$$\text{PastCost}(s) + h(s)$$

$$\leq$$

$$\text{PastCost}(s_{\text{end}})$$

# Amount explored




- If  $h(s) = 0$ , then  $A^*$  is same as UCS.
- If  $h(s) = \text{FutureCost}(s)$ , then  $A^*$  only explores nodes on a minimum cost path.
- Usually  $h(s)$  is somewhere in between.

# A\* search




**Key idea: distortion**

A\* distorts edge costs to favor end states.




# Admissibility



## Definition: admissibility

A heuristic  $h(s)$  is admissible if  
$$h(s) \leq \text{FutureCost}(s)$$

Intuition: admissible heuristics are optimistic



## Theorem: consistency implies admissibility

If a heuristic  $h(s)$  is **consistent**, then  $h(s)$  is **admissible**.

Proof: use induction on  $\text{FutureCost}(s)$



# Search: A\* relaxations





How do we get good heuristics? Just relax...



# Relaxation

**Intuition:** ideally, use  $h(s) = \text{FutureCost}(s)$ , but that's as hard as solving the original problem.



**Key idea: relaxation**

Constraints make life hard. Get rid of them.  
But this is just for the heuristic!





# Relaxation overview

