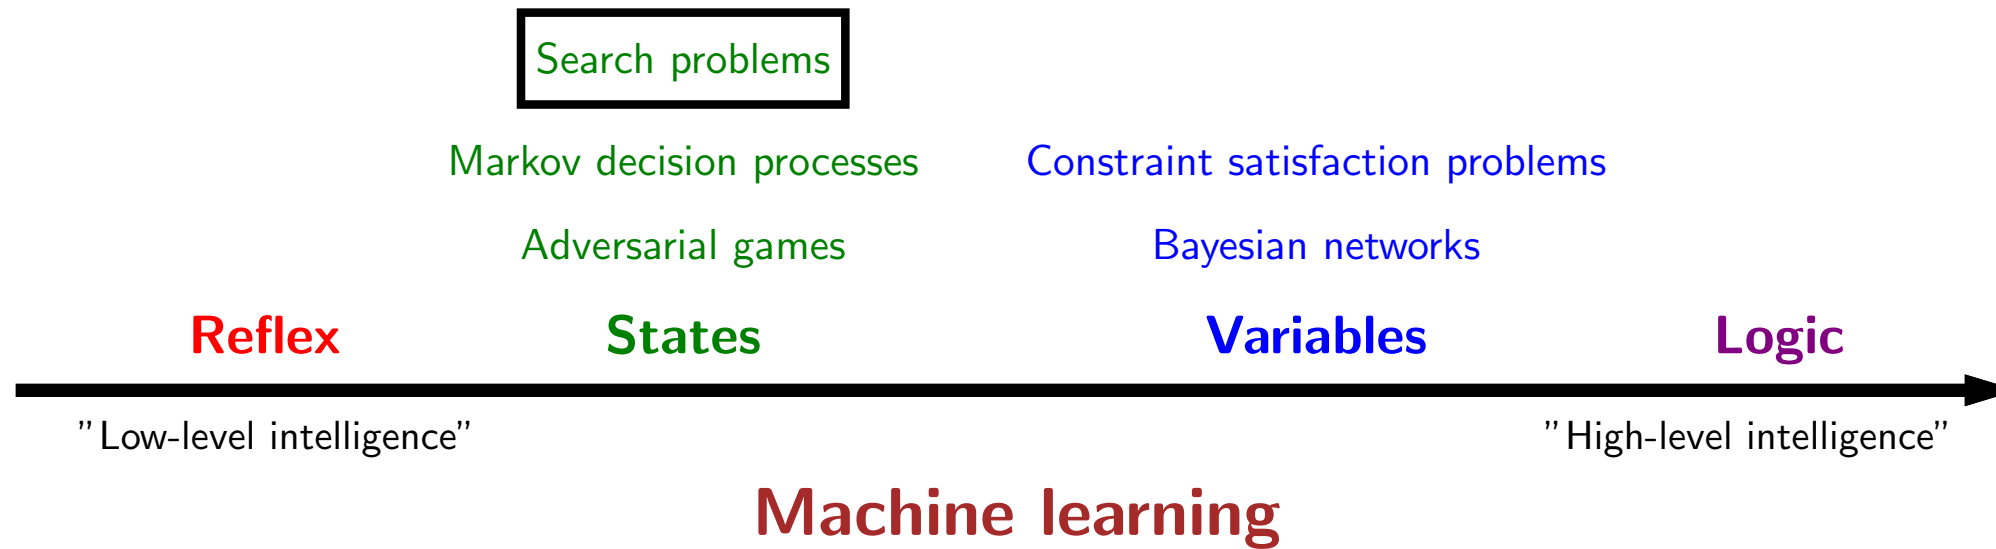





Search: overview



Course plan



Application: route finding



Objective: shortest? fastest? most scenic?

Actions: go straight, turn left, turn right

Application: robot motion planning



Objective: fastest path

Actions: acceleration and throttle

Application: robot motion planning



Objective: fastest? most energy efficient? safest? most expressive?

Actions: translate and rotate joints

Application: multi-robot systems



Objective: fastest? most energy efficient?

Actions: acceleration and steering of all robots

Application: machine translation

la maison bleue



the blue house

Objective: fluent English and preserves meaning

Actions: append single words (e.g., the)

Beyond reflex

Classifier (reflex-based models):



Search problem (state-based models):



Key: need to consider future consequences of an action!

Roadmap

Modeling

Modeling Search Problems

Algorithms

Tree Search

Dynamic Programming

Uniform Cost Search

Programming and Correctness of UCS

A*

A* Relaxations

Learning

Structured Perceptron

Paradigm

Modeling


Inference

Learning



Search: modeling






Farmer Cabbage Goat Wolf



Actions:

F▷	F◁
FC▷	FC◁
FG▷	FG◁
FW▷	FW◁

Approach: build a **search tree** ("what if?")



Search problem



Definition: search problem

- s_{start} : starting state
- $\text{Actions}(s)$: possible actions
- $\text{Cost}(s, a)$: action cost
- $\text{Succ}(s, a)$: successor
- $\text{IsEnd}(s)$: reached end state?



Transportation example



Example: transportation

Street with blocks numbered 1 to n .

Walking from s to $s + 1$ takes 1 minute.

Taking a magic tram from s to $2s$ takes 2 minutes.

How to travel from 1 to n in the least time?


[semi-live solution: `TransportationProblem`]



Search: tree search



Backtracking search



[whiteboard: search tree]

If b actions per state, maximum depth is D actions:

- **Memory:** $O(D)$ (small)
- **Time:** $O(b^D)$ (huge) [$2^{50} = 1125899906842624$]

Backtracking search



Algorithm: backtracking search

```
def backtrackingSearch( $s$ , path):  
    If IsEnd( $s$ ): update minimum cost path  
    For each action  $a \in \text{Actions}(s)$ :  
        Extend path with Succ( $s, a$ ) and Cost( $s, a$ )  
        Call backtrackingSearch(Succ( $s, a$ ), path)  
    Return minimum cost path
```

[semi-live solution: backtrackingSearch]

Depth-first search



Assumption: zero action costs

Assume action costs $\text{Cost}(s, a) = 0$.

Idea: Backtracking search + stop when find the first end state.

If b actions per state, maximum depth is D actions:

- **Space:** still $O(D)$
- **Time:** still $O(b^D)$ worst case, but could be much better if solutions are easy to find

Breadth-first search



Assumption: constant action costs

Assume action costs $\text{Cost}(s, a) = c$ for some $c \geq 0$.

Idea: explore all nodes in order of increasing depth.

Legend: b actions per state, solution has d actions

- **Space:** now $O(b^d)$ (a lot worse!)
- **Time:** $O(b^d)$ (better, depends on d , not D)

DFS with iterative deepening



Assumption: constant action costs

Assume action costs $\text{Cost}(s, a) = c$ for some $c \geq 0$.

Idea:

- Modify DFS to stop at a maximum depth.
- Call DFS for maximum depths $1, 2, \dots$

DFS on d asks: is there a solution with d actions?

Legend: b actions per state, solution size d

- **Space:** $O(d)$ (saved!)
- **Time:** $O(b^d)$ (same as BFS)



Tree search algorithms

Legend: b actions/state, solution depth d , maximum depth D

Algorithm	Action costs	Space	Time
Backtracking	any	$O(D)$	$O(b^D)$
DFS	zero	$O(D)$	$O(b^D)$
BFS	constant ≥ 0	$O(b^d)$	$O(b^d)$
DFS-ID	constant ≥ 0	$O(d)$	$O(b^d)$

- Always exponential time
- Avoid exponential space with DFS-ID

Tree Search Review

